Supervised Learning: K-Nearest Neighbors

- Define a distance metric
  - Euclidean
  - Manhatten
  - Any vector norm
- Choose the number of K neighbors
- Find the k nearest neighbors of the new observation that we want to classify

Prediction problem: what group is an observation associated with?

- Assign class label by majority vote
- Important to find the right K
  - Commonly use $k = \frac{\sqrt{N}}{2}$ where N = number of samples

K - number of neighbors that are closest to the new observation

K small: observation is local

K large: observation is an average of the neighborhood observations in the training data

More dims, need more data ↑

Curse of dimensionality - the more dimensions, the sparser in space the data points

Non-parametric, instance-based, lazy

- Non-parametric; Model is not defined by fixed set of parameters
- Instance-based or lazy learning: Model is the result of effectively memorizing training data
- Requires keeping the original data set
- Space and time complexity grow with size of training data
- Suffers from curse of dimensionality: points become increasingly isolated with more dims, for a fixed size training set
- sklearn.neighbors.KNeighborsClassifier